

# Docker Swarm:

## The Problem(s)

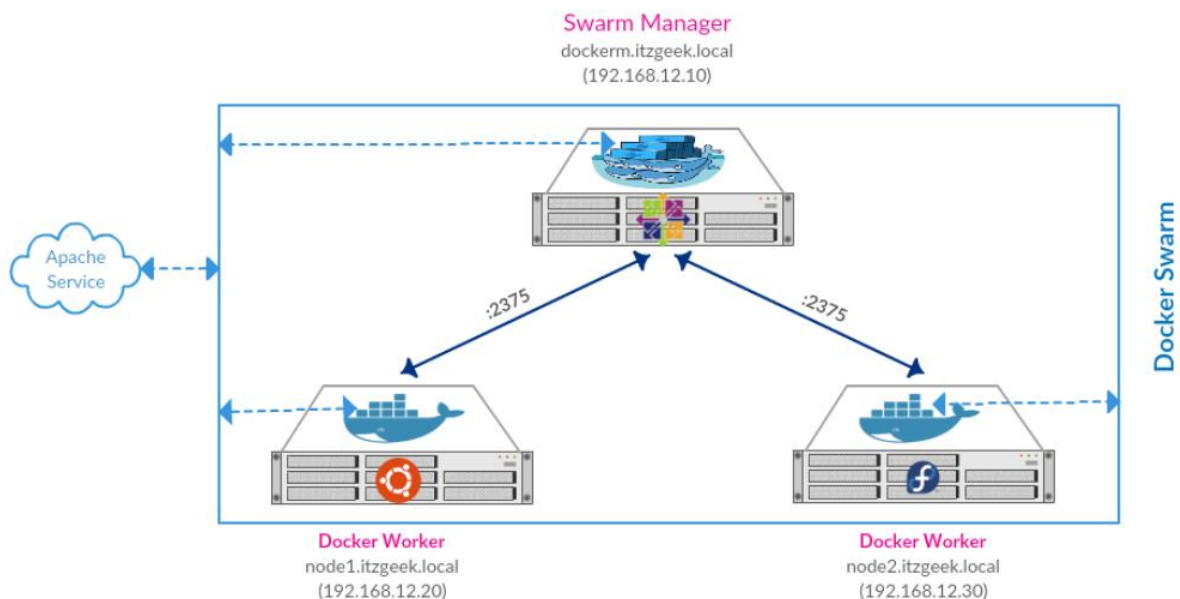
Container Orchestration:

Containers are great, but when you get lots of them running, at some point, you need them all working together to solve business problems.

The problem is, when you have lots of containers running, they need to be managed: you want enough of them running to handle the load, but not so many they are bogging down the machines in the cluster. And well, from time to time, containers are going to crash, and need to be restarted.

## Docker Swarm

Docker swarm mode allows you to manage a cluster of Docker Engines, natively within the Docker platform. You can use the Docker CLI to create a swarm, deploy application services to a swarm, and manage swarm behaviour.



- Coordinate between containers and allocate tasks to groups of containers
- Perform health checks and manage lifecycle of individual containers
- Provide redundancy and failover in case nodes experience failure
- Scale the number of containers up and down depending on load
- Perform rolling updates of software across multiple containers

## Lab :

1. Create one ec2 instance (called as manager), install docker

```
yum install -y docker
```

```
service docker start
```

2. Add node to docker swarm

```
docker swarm init
```

```
[root@ip-172-31-8-30 ec2-user]# docker swarm init
Swarm initialized: current node (jjaiubblyc3smu3rfja5uwo6b) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-45cj1l7ii0rxst1lvtgx23yik0emno2tizdhxw62d
oof9v32mr-4z1xxsclrnk5ibnu9snd7i0f4 172.31.8.30:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow
the instructions.
```

3. Create one more instance of ec2 (called as worker node), install docker on it.

4. Now make this instance as worker node to swarm.

5. copy the output of (docker swarm init) command and paste on worker node.

```
[root@ip-172-31-6-152 ec2-user]# docker swarm join --token SWMTKN-1-45cj1l7ii0rx
st1lvtgx23yik0emno2tizdhxw62doof9v32mr-4z1xxsclrnk5ibnu9snd7i0f4 172.31.8.30:237
7
This node joined a swarm as a worker.
```

6. Check if the node has been added to swarm, on manager run

```
docker node ls
```

```
[root@ip-172-31-8-30 ec2-user]# docker node ls
ID                                HOSTNAME                STATUS    AVAILABILITY    MANAGER STATUS
uxvh17am7c6hje6osvpye3iwr       ip-172-31-6-152        Ready    Active
jjaiubblyc3smu3rfja5uwo6b *     ip-172-31-8-30        Ready    Active           Leader
```

7. Check networks with

```
docker network ls
```

```
[root@ip-172-31-8-30 ec2-user]# docker network ls
NETWORK ID          NAME                DRIVER            SCOPE
aa66f0c0a651        bridge             bridge            local
26e942cac6d0        docker_gwbridge     bridge            local
60b44ff38da4        host               host              local
t8r4ufrkbped        ingress            overlay           swarm
99caee7cdf48        none              null              local
```

8. Each of them now has an overlay network called **ingress** and a bridge network called **docker\_gwbridge**

// The docker\_gwbridge connects the ingress network to the Docker host's network interface so that traffic can flow to and from swarm managers and workers. If you create swarm services and do not specify a network, they are connected to the ingress network. It is recommended that you use separate overlay networks for each application or group of applications which will work together. In the next procedure, you will create two overlay networks and connect a service to each of them.//

9. Now on manager node create an attachable overlay network called "Overlay-net"

docker network create --driver=overlay --attachable overlay-net

```
[root@ip-172-31-8-30 ec2-user]# docker network create --driver=overlay --attachable overlay-net
4545alumqdabb5xgk2zgy0xwk
```

// Notice the returned NETWORK ID -- you will see it again when you connect to it from worker node//

10. New network has been created only on master but not on node

11. On manager node create an interactive container alpine1 (Small. Simple. Secure. Alpine Linux is a security-oriented, lightweight Linux distribution based) that connects to overlay-net

docker run -it --name alpine1 --network overlay-net alpine

```
[root@ip-172-31-8-30 ec2-user]# docker run -it --name alpine1 --network overlay-net alpine
Unable to find image 'alpine:latest' locally
latest: Pulling from library/alpine
ff3a5c916c92: Pull complete
Digest: sha256:7df6db5aa61ae9480f52f0b3a06a140ab98d427f86d8d5de0bedab9b8df6b1c0
Status: Downloaded newer image for alpine:latest
/ #
```

12. Check on worker node, but the overlay-net is still not present (docker network ls)

13. On worker node create an alpine2 container in interactive & detached mode.

docker run -dit --name alpine2 --network overlay-net alpine

```
[root@ip-172-31-6-152 ec2-user]# docker run -dit --name alpine2 --network overlay-net alpine
Unable to find image 'alpine:latest' locally
latest: Pulling from library/alpine
ff3a5c916c92: Pull complete
Digest: sha256:7df6db5aa61ae9480f52f0b3a06a140ab98d427f86d8d5de0bedab9b8df6b1c0
Status: Downloaded newer image for alpine:latest
13d939057a1f5a55d3aa1b63a3ccb446bec40fcc624fb50c1d44d2a59ea9db9b
```

// Automatic DNS container discovery only works with unique container names//

14. Now overlay-net will be listed on worker node

15. Try to ping container alpine2 running (container must be running, that's why we have started it in detached mode) on worker node from alpine1 running on manager node.

ping -c 2 alpine2

```
/ # ping -c 2 alpine2
PING alpine2 (10.0.0.5): 56 data bytes
64 bytes from 10.0.0.5: seq=0 ttl=255 time=0.535 ms
64 bytes from 10.0.0.5: seq=1 ttl=255 time=0.688 ms

--- alpine2 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 0.535/0.611/0.688 ms
/ #
```

16. Ping alpine1 from alpine2 and should be able to reach.

```
/ # ping -c 2 alpine1
PING alpine1 (10.0.0.6): 56 data bytes
64 bytes from 10.0.0.6: seq=0 ttl=255 time=0.611 ms
64 bytes from 10.0.0.6: seq=1 ttl=255 time=0.693 ms

--- alpine1 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 0.611/0.652/0.693 ms
```